

Introduction to Database Systems

CSE 414

Lecture 5: SQL Aggregates

`Product(pname, price, category, manufacturer)`

`Company(cname, country)`

Joins in SQL

pname	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all Japanese products that cost < \$150

```
SELECT P.pname, P.price  
FROM Product as P, Company as C  
WHERE P.manufacturer=C.cname AND  
C.country='Japan' AND C.price < 150
```

Join Predicate

Selection predicates

(Inner) Joins

```
SELECT x1.a1, x2.a2, ... xm.am  
FROM   R1 as x1, R2 as x2, ... Rm as xm  
WHERE  Cond
```

✓ tuple in R1

for x1 in R1:

 for x2 in R2:

 ...

 for xm in Rm:

 if Cond(x1, x2...):

 output(x1.a1, x2.a2, ... xm.am)

This is called nested loop semantics since we are interpreting what a join means using a nested loop

Self Join Example

```
Product(pname, price, category, manufacturer)  
Company(cname, country)
```

-- manufacturer is foreign key to Company

Find US companies that manufacture both
'gadgets' and 'photo' products

```
SELECT DISTINCT z.cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
AND x.manufacturer = z.cname  
AND y.manufacturer = z.cname  
AND x.category = 'gadget'  
AND y.category = 'photography';
```

Need to include
Product twice!

Joins in SQL

- The join we have just seen is sometimes called an **inner join**
 - Each row in the result **must come from both tables in the join**
- Sometimes we want to include rows from only one of the two table: **outer join**

Employee(id, name)

Sales(employeeID, productID)

Outer Join

left

right

Employee

<u>id</u>	name
1	Joe
2	Jack
3	Jill

Sales

<u>employeeID</u>	productID
1	344
1	355
2	544

Retrieve employees and their sales

Jill is present

```
SELECT *
FROM Employee E
LEFT OUTER JOIN
Sales S
ON E.id = S.employeeID
```

<u>id</u>	name	<u>employeeID</u>	productID
1	Joe	1	344
1	Joe	1	355
2	Jack	2	544
3	Jill	NULL	NULL

Outer Joins

```
tableA (LEFT/RIGHT/FULL) OUTER JOIN tableB ON p
```

- Left outer join:
 - Include tuples from tableA even if no match
- Right outer join:
 - Include tuples from tableB even if no match
- Full outer join:
 - Include tuples from both even if no match
- In all cases:
 - Patch tuples without matches using NULL

Aggregates in SQL

Simple Aggregations

Five basic aggregate operations in SQL

```
select count(*) from Purchase  
select sum(quantity) from Purchase  
select avg(price) from Purchase  
select max(quantity) from Purchase  
select min(quantity) from Purchase
```

Except count, all aggregations apply to a single attribute

Demo

Aggregates and NULL Values

Null values are not used in aggregates

```
insert into Purchase  
values(12, 'gadget', NULL, NULL, 'april')
```

Let's try the following

```
select count(*) from Purchase
```

```
select count(quantity) from Purchase
```

```
select sum(quantity) from Purchase
```

```
select count(*)
```

```
from Purchase
```

```
where quantity is not null;
```

Counting Duplicates

COUNT applies to duplicates, unless otherwise stated:

```
SELECT count(product)
FROM Purchase
WHERE price > 4.99
```

same as count(*) if no nulls

We probably want:

```
SELECT count(DISTINCT product)
FROM Purchase
WHERE price > 4.99
```

More Examples

```
SELECT Sum(P.price * P.quantity)  
FROM Purchase as P
```

What do
they mean ?

```
SELECT Sum(P.price * P.quantity)  
FROM Purchase as P  
WHERE P.product = 'bagel'
```


Grouping and Aggregation

Purchase(product, price, quantity)

Find total quantities for all sales over \$1, by product.

Grouping and Aggregation

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10



Product	TotalSales
Bagel	40
Banana	20

```
SELECT      product, Sum(quantity) AS TotalSales  
FROM        Purchase  
WHERE       price > 1  
GROUP BY    product
```

Other Examples

Compare these
two queries:

```
SELECT product, count(*)  
FROM Purchase  
GROUP BY product
```

```
SELECT month, count(*)  
FROM Purchase  
GROUP BY month
```

```
SELECT product,  
       sum(quantity) AS SumQuantity,  
       max(price) AS MaxPrice  
FROM Purchase  
GROUP BY product
```

What does
it return?

Need to be Careful...

```
SELECT product,  
       max(quantity)  
FROM   Purchase  
GROUP  BY product
```

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Need to be Careful...

```
SELECT product,  
       max(quantity)  
FROM Purchase  
GROUP BY product
```

```
SELECT product, quantity  
FROM Purchase  
GROUP BY product  
-- what does this mean?
```

Product	Price	Quantity
Bagel	3	20
	1.50	20
Banana	0.5	50
	2	10
	4	10

Need to be Careful...

```
SELECT product,  
       max(quantity)  
FROM Purchase  
GROUP BY product
```

```
SELECT product, quantity  
FROM Purchase  
GROUP BY product  
-- what does this mean?
```

Product	Price	Quantity
Bagel	3	20
	1.50	20
Banana	0.5	50
	2	10
	4	10

Product	Max(quantity)
Bagel	20
Banana	50

INF

Need to be Careful...

```
SELECT product,  
       max(quantity)  
FROM   Purchase  
GROUP  BY product
```

```
SELECT      product, quantity  
FROM        Purchase  
GROUP BY   product  
-- what does this mean?
```

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

Product	Quantity
Bagel	20
Banana	??

Need to be Careful...

```
SELECT product,  
       max(quantity)  
FROM   Purchase  
GROUP BY product
```

```
SELECT      product, quantity  
FROM        Purchase  
GROUP BY   product  
-- what does this mean?
```

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

Product	Quantity
Bagel	20
Banana	??



Everything in SELECT must be either a GROUP-BY attribute, or an aggregate

Need to be Careful...

```
SELECT product,  
       max(quantity)  
FROM Purchase  
GROUP BY product
```

```
SELECT product, quantity  
FROM Purchase  
GROUP BY product  
-- what does this mean?
```

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

Product	Quantity
Bagel	20
Banana	??



Grouping and Aggregation

Purchase(product, price, quantity)

Find total quantities for all sales over \$1, by product.

```
SELECT      product, Sum(quantity) AS TotalSales  
FROM        Purchase  
WHERE       price > 1  
GROUP BY    product
```

How is this query processed?

Grouping and Aggregation

Purchase(product, price, quantity)

Find total quantities for all sales over \$1, by product.

```
SELECT      product, Sum(quantity) AS TotalSales  
FROM        Purchase  
WHERE       price > 1  
GROUP BY    product
```

Do these queries return the same number of rows? Why?

```
SELECT      product, Sum(quantity) AS TotalSales  
FROM        Purchase  
GROUP BY    product
```

Grouping and Aggregation

Purchase(product, price, quantity)

Find total quantities for all sales over \$1, by product.

```
SELECT      product, Sum(quantity) AS TotalSales  
FROM        Purchase  
WHERE       price > 1  
GROUP BY    product
```

Do these queries return the same number of rows? Why?

```
SELECT      product, Sum(quantity) AS TotalSales  
FROM        Purchase  
GROUP BY    product
```

Rows where price > 1 are removed, so first query may return fewer groups

Grouping and Aggregation

1. Compute the **FROM** and **WHERE** clauses.
2. Group by the attributes in the **GROUPBY**
3. Compute the **SELECT** clause:
grouped attributes and aggregates.

The logo consists of the letters "FWGS" in a bold, red, sans-serif font, enclosed within a thin red rectangular border.

TM

1,2: From, Where

FWGS

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10


WHERE price > 1

```
SELECT product, Sum(quantity) AS TotalSales  
FROM Purchase  
WHERE price > 1  
GROUP BY product
```

3,4. Grouping, Select

FWGS

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10



Product	TotalSales
Bagel	40
Banana	20

```
SELECT product, Sum(quantity) AS TotalSales  
FROM Purchase  
WHERE price > 1  
GROUP BY product
```

Purchase(pid, product, price, quantity, month)

Ordering Results

```
SELECT product, sum(price*quantity) as rev  
FROM Purchase  
GROUP BY product  
ORDER BY rev desc
```

FWGOS

TM

Note: some SQL engines
want you to say ORDER BY sum(price*quantity) desc

Purchase(pid, product, price, quantity, month)

HAVING Clause

Same query as before, except that we consider only products that had at least 30 sales.

```
SELECT      product, sum(price*quantity)
FROM        Purchase
WHERE       price > 1
GROUP BY    product
HAVING     sum(quantity) > 30
```

HAVING clause contains conditions on aggregates.

General form of Grouping and Aggregation

```
SELECT      S
FROM        R1, ..., Rn
WHERE       C1
GROUP BY   a1, ..., ak
HAVING     C2
```

Why ?

S = may contain attributes a_1, \dots, a_k and/or any aggregates but NO OTHER ATTRIBUTES

C₁ = is any condition on the attributes in R_1, \dots, R_n

C₂ = is any condition on aggregate expressions and on attributes a_1, \dots, a_k

Semantics of SQL With Group-By

```
SELECT      S  
FROM        R1, ..., Rn  
WHERE       C1  
GROUP BY   a1, ..., ak  
HAVING     C2
```

FWGHOS

Evaluation steps:

1. Evaluate FROM-WHERE using Nested Loop Semantics
2. Group by the attributes a_1, \dots, a_k
3. Apply condition C_2 to each group (may have aggregates)
4. Compute aggregates in S and return the result