

Introduction to Database Systems

CSE 414

Lecture 6: SQL Subqueries

Announcements

- HW2 and WQ2 released
 - Both due next Tuesday
- Please fill in the Azure questionnaire by tonight!
 - See HW2 writeup for details

Simple Aggregations

Five basic aggregate operations in SQL

```
select count(*) from Purchase  
select sum(quantity) from Purchase  
select avg(price) from Purchase  
select max(quantity) from Purchase  
select min(quantity) from Purchase
```

Except count, all aggregations apply to a single attribute

Everything in SELECT must be either a GROUP-BY attribute, or an aggregate

Need to be Careful...

```
SELECT product,  
       max(quantity)  
FROM Purchase  
GROUP BY product
```

```
SELECT product, quantity  
FROM Purchase  
GROUP BY product  
-- what does this mean?
```

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

Product	Quantity
Bagel	20
Banana	??

INF



Semantics of SQL With Group-By

```
SELECT      S  
FROM        R1, ..., Rn  
WHERE       C1  
GROUP BY   a1, ..., ak  
HAVING     C2
```

FWGHOS

Evaluation steps:

1. Evaluate FROM-WHERE using Nested Loop Semantics
2. Group by the attributes a_1, \dots, a_k
3. Apply condition C_2 to each group (may have aggregates)
4. Compute aggregates in S and return the result

Purchase(pid, product, price, quantity, month)

Exercise

FWGHOS

Compute the total income per month

Show only months with less than 10 items sold

Order by quantity sold and display as “TotalSold”

Purchase(pid, product, price, quantity, month)

Exercise

FWGHOS

Compute the total income per month

Show only months with less than 10 items sold

Order by quantity sold and display as “TotalSold”

FROM

Purchase

Purchase(pid, product, price, quantity, month)

Exercise

FWGHOS

Compute the total income per month

Show only months with less than 10 items sold

Order by quantity sold and display as “TotalSold”

FROM Purchase
GROUP BY month

Purchase(pid, product, price, quantity, month)

Exercise

FWGHOS

Compute the total income per month

Show only months with less than 10 items sold

Order by quantity sold and display as “TotalSold”

```
FROM Purchase
GROUP BY month
HAVING sum(quantity) < 10
```

Purchase(pid, product, price, quantity, month)

Exercise

FWGHOS

Compute the total income per month

Show only months with less than 10 items sold

Order by quantity sold and display as “TotalSold”

```
SELECT      month, sum(price*quantity),  
            sum(quantity) as TotalSold  
FROM        Purchase  
GROUP BY    month  
HAVING      sum(quantity) < 10
```

Purchase(pid, product, price, quantity, month)

Exercise

FWGHOS

Compute the total income per month

Show only months with less than 10 items sold

Order by quantity sold and display as “TotalSold”

```
SELECT      month, sum(price*quantity),  
            sum(quantity) as TotalSold  
FROM        Purchase  
GROUP BY    month  
HAVING      sum(quantity) < 10  
ORDER BY    sum(quantity)
```

WHERE vs HAVING

- WHERE condition is applied to individual rows
 - The rows may or may not contribute to the aggregate
 - No aggregates allowed here
- HAVING condition is applied to the entire group
 - Only applicable if GROUP BY is involved
 - Entire group is returned, or removed
 - May use aggregate functions on the group

`Product(pid, pname, manufacturer)`

`Purchase(id, product_id, price, month)`

Aggregate + Join

For each manufacturer, compute how many products
with price > \$100 they sold

`Product(pid, pname, manufacturer)`

`Purchase(id, product_id, price, month)`

Aggregate + Join

For each manufacturer, compute how many products with price > \$100 they sold

Problem: manufacturer is in Product, price is in Purchase...

Product(pid,pname,manufacturer)

Purchase(id,product_id,price,month)

Aggregate + Join

For each manufacturer, compute how many products with price > \$100 they sold

Problem: manufacturer is in Product, price is in Purchase...

```
-- step 1: think about their join  
SELECT ...  
FROM Product x, Purchase y  
WHERE x.pid = y.product_id  
and y.price > 100
```

manu facturer	...	price	...
Hitachi		150	
Canon		300	
Hitachi		180	

Product(pid,pname,manufacturer)

Purchase(id,product_id,price,month)

Aggregate + Join

For each manufacturer, compute how many products with price > \$100 they sold

Problem: manufacturer is in Product, price is in Purchase...

```
-- step 1: think about their join  
SELECT ...  
FROM Product x, Purchase y  
WHERE x.pid = y.product_id  
and y.price > 100
```

manu facturer	...	price	...
Hitachi		150	
Canon		300	
Hitachi		180	

```
-- step 2: do the group-by on the join  
SELECT x.manufacturer, count(*)  
FROM Product x, Purchase y  
WHERE x.pid = y.product_id  
and y.price > 100  
GROUP BY x.manufacturer
```

manu facturer	count(*)
Hitachi	2
Canon	1
...	

Product(pid,pname,manufacturer)

Purchase(id,product_id,price,month)

Aggregate + Join

Variant:

For each manufacturer, compute how many products with price > \$100 they sold **in each month**

```
SELECT x.manufacturer, y.month, count(*)  
FROM Product x, Purchase y  
WHERE x.pid = y.product_id  
and y.price > 100  
GROUP BY x.manufacturer, y.month
```

manu facturer	month	count(*)
Hitachi	Jan	2
Hitachi	Feb	1
Canon	Jan	3
...		

Including Empty Groups


- In the result of a group by query, there is one row per group in the result

Count(*) is never 0

```
SELECT x.manufacturer, count(*)  
FROM Product x, Purchase y  
WHERE x.pname = y.product  
GROUP BY x.manufacturer
```

Including Empty Groups

```
SELECT x.manufacturer, count(y.pid)
FROM Product x LEFT OUTER JOIN Purchase y
ON x.pname = y.product
GROUP BY x.manufacturer
```



Count(pid) is 0
when all pid's in
the group are
NULL

What we have in our SQL toolbox

- Projections (SELECT * / SELECT c1, c2, ...)
- Selections (aka filtering) (WHERE cond)
- Joins (inner and outer)
- Aggregates
- Group by
- Inserts, updates, and deletes

Make sure you read the textbook!

Subqueries

- A subquery is a SQL query nested inside a larger query
- Such inner-outer queries are called nested queries
- A subquery may occur in:
 - A SELECT clause
 - A FROM clause
 - A WHERE clause
- **Rule of thumb: avoid nested queries when possible**
 - But sometimes it's impossible, as we will see

FWGHOS

Subqueries...

- Can return a single value to be included in a SELECT clause
- Can return a relation to be included in the FROM clause, aliased using a tuple variable
- Can return a single value to be compared with another value in a WHERE clause
- Can return a relation to be used in the WHERE or HAVING clause under an existential quantifier

1. Subqueries in SELECT

Product (pname, price, cid)

Company (cid, cname, city)

For each product return the city where it is manufactured

```
SELECT X.pname, (SELECT Y.city  
                  FROM Company Y  
                 WHERE Y.cid=X.cid) as City  
FROM Product X
```

“correlated
subquery”

What happens if the subquery returns more than one city?

We get a runtime error

(and SQLite simply ignores the extra values...)

Product (pname, price, cid)

Company (cid, cname, city)

1. Subqueries in SELECT

Whenever possible, don't use a nested queries:

```
SELECT X.pname, (SELECT Y.city  
                  FROM Company Y  
                 WHERE Y.cid=X.cid) as City  
FROM Product X
```

||

```
SELECT X.pname, Y.city  
FROM Product X, Company Y  
WHERE X.cid=Y.cid
```

We have
“unnested”
the query

Product (pname, price, cid)

Company (cid, cname, city)

1. Subqueries in SELECT

Compute the number of products made by each company

```
SELECT DISTINCT C.cname, (SELECT count(*)  
                           FROM Product P  
                          WHERE P.cid=C.cid)  
FROM Company C
```

Product (pname, price, cid)

Company (cid, cname, city)

1. Subqueries in SELECT

Compute the number of products made by each company

```
SELECT DISTINCT C.cname, (SELECT count(*)  
                           FROM Product P  
                          WHERE P.cid=C.cid)  
      FROM Company C
```

Better: we can
unnest using a
GROUP BY

```
SELECT C.cname, count(*)  
      FROM Company C, Product P  
     WHERE C.cid=P.cid  
   GROUP BY C.cname
```

Product (pname, price, cid)

Company (cid, cname, city)

1. Subqueries in SELECT

But are these really equivalent?

```
SELECT DISTINCT C.cname, (SELECT count(*)  
                           FROM Product P  
                          WHERE P.cid=C.cid)  
      FROM Company C
```

```
SELECT C.cname, count(*)  
  FROM Company C, Product P  
 WHERE C.cid=P.cid  
 GROUP BY C.cname
```

Product (pname, price, cid)

Company (cid, cname, city)

1. Subqueries in SELECT

But are these really equivalent?

```
SELECT DISTINCT C.cname, (SELECT count(*)  
                           FROM Product P  
                          WHERE P.cid=C.cid)  
      FROM Company C
```

```
SELECT C.cname, count(*)  
  FROM Company C, Product P  
 WHERE C.cid=P.cid  
 GROUP BY C.cname
```

No! Different results if a company has no products

```
SELECT C.cname, count(pname)  
  FROM Company C LEFT OUTER JOIN Product P  
    ON C.cid=P.cid  
 GROUP BY C.cname
```

Product (pname, price, cid)

Company (cid, cname, city)

2. Subqueries in FROM

Find all products whose prices is > 20 and < 500

```
SELECT X.pname
FROM (SELECT *
      FROM Product AS Y
      WHERE price > 20) as X
WHERE X.price < 500
```

Product (pname, price, cid)

Company (cid, cname, city)

2. Subqueries in FROM

Find all products whose prices is > 20 and < 500

```
SELECT X.pname
FROM (SELECT *
      FROM Product AS Y
      WHERE price > 20) as X
WHERE X.price < 500
```

Try unnest this query !

Product (pname, price, cid)

Company (cid, cname, city)

2. Subqueries in FROM

Find all products whose prices is > 20 and < 500

```
SELECT X.pname
FROM (SELECT *
      FROM Product AS Y
      WHERE price > 20) as X
WHERE X.price < 500
```

Side note: This is not a correlated subquery. (why?)

Try unnest this query !

2. Subqueries in FROM

Sometimes we need to compute an intermediate table only to use it later in a SELECT-FROM-WHERE

- Option 1: use a subquery in the FROM clause
- Option 2: use the WITH clause
 - See textbook for details

Product (pname, price, cid)

Company (cid, cname, city)

2. Subqueries in FROM

```
SELECT X.pname  
FROM (SELECT *  
      FROM Product AS Y  
      WHERE price > 20) as X  
WHERE X.price < 500
```

||

A subquery whose
result we called myTable

```
WITH myTable AS (SELECT * FROM Product AS Y WHERE price > 20)  
SELECT X.pname  
FROM myTable as X  
WHERE X.price < 500
```