

Introduction to Database Systems CSE 414

Lecture 8: Datalog

CSE 414 - Spring 2018

1

Announcements

- HW3 posted (1 week)
 - Same dataset, more challenging queries
 - We have sent out all Azure codes if you filled out the form earlier
 - Make sure you use the cheapest tier
 - aka **READ THE HW INSTRUCTIONS**
 - You should first run on sqlite in any case!

CSE 414 - Spring 2018

2

Class Overview

- Unit 1: Intro
- Unit 2: Relational Data Models and Query Languages
 - Data models, SQL, **Datalog**, Relational Algebra
- Unit 3: Non-relational data
- Unit 4: RDBMS internals and query optimization
- Unit 5: Parallel query processing
- Unit 6: DBMS usability, conceptual design
- Unit 7: Transactions

CSE 414 - Spring 2018

3

What is Datalog?

- Another query language for relational model
 - Designed in the 80's
 - Simple, concise, elegant
 - Extends relational queries with recursion
- Today is a hot topic:
 - Souffle (we will use in HW4)
 - Eve <http://witheve.com/>
 - Differential datalog
<https://github.com/frankmcsherry/differential-dataflow>
 - Beyond databases in many research projects: network protocols, static program analysis

4



- Open-source implementation of Datalog DBMS
- Under active development
- Commercial implementations are available
 - More difficult to set up and use
- “sqlite” of Datalog
 - Set-based rather than bag-based
- Install in your VM
 - Run `sudo yum install souffle` in terminal
 - More details in upcoming HW4

5

Why bother with *yet* another
relational query language?

CSE 414 - Spring 2018

6

Example: storing FB friends

As a graph

Or

Person1	Person2	is_friend
Peter	John	1
John	Mary	0
Mary	Phil	1
Phil	Peter	1
...

As a relation

We will learn the tradeoffs of different data models later this quarter

CSE 414 - Spring 2018 7

Compute your friends graph

```
SELECT f.p2
FROM Friends as f
WHERE f.p1 = 'me' AND f.isFriend = 1
```

My own friends

```
SELECT f1.p2
FROM Friends as f1,
     (SELECT f.p2
      FROM Friends as f
      WHERE f.p1 = 'me' AND
            f.isFriend = 1) as f2
WHERE f1.p1 = f2.p2 AND
      f1.isFriend = 1
```

My FoF

My FoFoF... My FoFoFoF..

When does it end???

p1	p2	isFriend
Peter	John	1
John	Mary	0
Mary	Phil	1
Phil	Peter	1
...

Friends(p1, p2, isFriend)

Datalog allows us to write recursive queries easily

CSE 414 - Spring 2018 7

Actor(id, fname, lname)
 Casts(pid, mid)
 Movie(id, name, year)

Schema

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

CSE 414 - Spring 2018 9

Actor(id, fname, lname)
 Casts(pid, mid)
 Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```
.decl Actor(id:number, fname:symbol, lname:symbol)
.decl Casts(id:number, mid:number)
.decl Movie(id:number, name:symbol, year:number)

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 29000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
```

Table declaration

Types in Souffle:
 number
 symbol (aka varchar)

Insert data

CSE 414 - Spring 2018 10

Actor(id, fname, lname)
 Casts(pid, mid)
 Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```
Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 29000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
```

```
Q1(y) :- Movie(x,y,z), z=1940.
```

CSE 414 - Spring 2018 11

Actor(id, fname, lname)
 Casts(pid, mid)
 Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```
Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 29000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
```

```
Q1(y) :- Movie(x,y,z), z=1940.
```

Find Movies made in 1940

CSE 414 - Spring 2018 12

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

SQL

```

SELECT name
FROM Movie
WHERE year = 1940
  
```

Find Movies made in 1940

CSE 414 - Spring 2018 13

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

Q1(y) :- Movie(x,y,z), z=1940.

Order of variable matters!

Find Movies made in 1940

CSE 414 - Spring 2018 14

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

Q1(y) :- Movie(iDontCare,y,z), z=1940.

Find Movies made in 1940

CSE 414 - Spring 2018 15

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

Q1(y) :- Movie(_,y,z), z=1940.

_ = "don't care" variables

Find Movies made in 1940

CSE 414 - Spring 2018 16

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

Q1(y) :- Movie(x,y,z), z=1940.

Q2(f,l) :- Actor(z,f,l), Casts(z,x), Movie(x,y,1940).

Find Actors who acted in Movies made in 1940

CSE 414 - Spring 2018 17

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

Q1(y) :- Movie(x,y,z), z=1940.

Q2(f,l) :- Actor(z,f,l), Casts(z,x), Movie(x,y,1940).

Find Actors who acted in Movies made in 1940

CSE 414 - Spring 2018 18

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

```

Q1(y) :- Movie(x,y,z), z=1940.
Q2(f,l) :- Actor(z,f,l), Casts(z,x),
           Movie(x,y,1940).
Q3(f,l) :- Actor(z,f,l), Casts(z,x1), Movie(x1,y1,1910),
           Casts(z,x2), Movie(x2,y2,1940).
  
```

CSE 414 - Spring 2018 19

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

```

Q1(y) :- Movie(x,y,z), z=1940.
Q2(f,l) :- Actor(z,f,l), Casts(z,x),
           Movie(x,y,1940).
Q3(f,l) :- Actor(z,f,l), Casts(z,x1), Movie(x1,y1,1910),
           Casts(z,x2), Movie(x2,y2,1940).
  
```

both

Find Actors who acted in a Movie in 1940 and in one in 1910

CSE 414 - Spring 2018 20

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Datalog: Facts and Rules

Facts = tuples in the database **Rules** = queries

```

Actor(344759, 'Douglas', 'Fowley').
Casts(344759, 29851).
Casts(355713, 298000).
Movie(7909, 'A Night in Armour', 1910).
Movie(29000, 'Arizona', 1940).
Movie(29445, 'Ave Maria', 1940).
  
```

```

Q1(y) :- Movie(x,y,z), z=1940.
Q2(f,l) :- Actor(z,f,l), Casts(z,x),
           Movie(x,y,1940).
Q3(f,l) :- Actor(z,f,l), Casts(z,x1), Movie(x1,y1,1910),
           Casts(z,x2), Movie(x2,y2,1940).
  
```

Extensional Database Predicates = EDB = Actor, Casts, Movie
Intensional Database Predicates = IDB = Q1, Q2, Q3

CSE 414 - Spring 2018 21

Datalog: Terminology

head

body

atom

atom

atom (aka subgoal)

```

Q2(f, l) :- Actor(z,f,l), Casts(z,x), Movie(x,y,1940).
  
```

f, l = head variables
x,y,z = existential variables

CSE 414 - Spring 2018 22

More Datalog Terminology

```

Q(args) :- R1(args), R2(args), ...
  
```

- $R_i(args_i)$ called an *atom*, or a *relational predicate*
- $R_i(args_i)$ evaluates to true when relation R_i contains the tuple described by $args_i$.
 - Example: Actor(344759, 'Douglas', 'Fowley') is true
- In addition we can also have arithmetic predicates
 - Example: $z > 1940$.
- Book uses AND instead of ,

```
Q(args) :- R1(args) AND R2(args) ...
```

Datalog program

- A Datalog program consists of several rules
- Importantly, rules may be recursive!
 - Recall CSE 143!
- Usually there is one distinguished predicate that's the output
- We will show an example first, then give the general semantics.

CSE 414 - Spring 2018 24

4

R encodes a graph
e.g., connected cities

Example

R=

1	2
2	1
2	3
1	4
3	4
4	5

R encodes a graph
e.g., connected cities

Example

R=

1	2
2	1
2	3
1	4
3	4
4	5

$T(x,y) :- R(x,y).$
 $T(x,y) :- R(x,z), T(z,y).$

Multiple rules for the same IDB means OR

What does it compute?

R encodes a graph
e.g., connected cities

Example

R=

1	2
2	1
2	3
1	4
3	4
4	5

Initially:
T is empty.

$T(x,y) :- R(x,y).$
 $T(x,y) :- R(x,z), T(z,y).$

What does it compute?

R encodes a graph
e.g., connected cities

Example

R=

1	2
2	1
2	3
1	4
3	4
4	5

Initially:
T is empty.

$T(x,y) :- R(x,y).$
 $T(x,y) :- R(x,z), T(z,y).$

First iteration:
T =

1	2
2	1
2	3
1	4
3	4
4	5

First rule generates this

Second rule generates nothing (because T is empty)

What does it compute?

R encodes a graph
e.g., connected cities

Example

R=

1	2
2	1
2	3
1	4
3	4
4	5

Initially:
T is empty.

First iteration:
T =

1	2
2	1
2	3
1	4
3	4
4	5

Second iteration:
T =

1	2
2	1
2	3
1	4
3	4
4	5
1	1
2	2
1	3
2	4
1	5
3	5

First rule generates this

Second rule generates this

New facts

$T(x,y) :- R(x,y).$
 $T(x,y) :- R(x,z), T(z,y).$

What does it compute?

R encodes a graph
e.g., connected cities

Example

R=

1	2
2	1
2	3
1	4
3	4
4	5

Initially:
T is empty.

First iteration:
T =

1	2
2	1
2	3
1	4
3	4
4	5

Second iteration:
T =

1	2
2	1
2	3
1	4
3	4
4	5
1	1
2	2
1	3
2	4
1	5
3	5

Third iteration:
T =

1	2
2	1
2	3
1	4
3	4
4	5
1	1
2	2
1	3
2	4
1	5
3	5
2	5

Both rules

First rule

Second rule

New fact

$T(x,y) :- R(x,y).$
 $T(x,y) :- R(x,z), T(z,y).$

What does it compute?

R encodes a graph
e.g., connected cities

Example

$T(x,y) :- R(x,y).$
 $T(x,y) :- R(x,z), T(z,y).$

What does it compute?

Initially: T is empty.

First iteration: T =

1	2
2	1
2	3
1	4
3	4
4	5

Second iteration: T =

1	2
2	1
2	3
1	4
3	4
4	5
1	1
2	2
1	3
2	4
1	5
3	5

Third iteration: T =

1	2
2	1
2	3
1	4
3	4
4	5
1	1
2	2
1	3
2	4
1	5
3	5

Fourth iteration: T = (same)
No new facts. DONE

R =

1	2
2	1
2	3
1	4
3	4
4	5

Datalog Semantics

Fixpoint semantics

- Start:
 - $IDB_0 = \text{empty relations}$
 - $t = 0$
- Repeat:
 - $IDB_{t+1} = \text{Compute Rules}(EDB, IDB_t)$
 - $t = t+1$
- Until $IDB_t = IDB_{t-1}$

Remark: since rules are **monotone**:
 $\emptyset = IDB_0 \subseteq IDB_1 \subseteq IDB_2 \subseteq \dots$

It follows that a datalog program w/o functions (+, *, ...) always terminates. (Why?)

Three Equivalent Programs

R encodes a graph
e.g., connected cities

R =

1	2
2	1
2	3
1	4
3	4
4	5

Right linear: $T(x,y) :- R(x,y).$
 $T(x,y) :- R(x,z), T(z,y).$

Left linear: $T(x,y) :- R(x,y).$
 $T(x,y) :- T(x,z), R(z,y).$

Non-linear: $T(x,y) :- R(x,y).$
 $T(x,y) :- T(x,z), T(z,y).$

Question: which terminates in fewest iterations?

CSE 414 - Spring 2018 33

More Features

- Aggregates
- Grouping
- Negation

CSE 414 - Spring 2018 34

Actor(id, fname, lname)
 Casts(pid, mid)
 Movie(id, name, year)

Aggregates

[aggregate name] <var> : { [relation to compute aggregate on] }

$\min x : \{ \text{Actor}(x, y, _), y = \text{'John'} \}$

$Q(\minId) :- \minId = \min x : \{ \text{Actor}(x, y, _), y = \text{'John'} \}$

Assign variable to the value of the aggregate

Meaning (in SQL)

```
SELECT min(id) as minId
FROM Actor as a
WHERE a.name = 'John'
```

Aggregates in Souffle:

- Count
- Min
- Max
- Sum

CSE 414 - Spring 2018 35

Actor(id, fname, lname)
 Casts(pid, mid)
 Movie(id, name, year)

Aggregates

[aggregate name] <var> : { [relation to compute aggregate on] }

$\min x : \{ \text{Actor}(x, y, _), y = \text{'John'} \}$

$Q(\minId, y) :- \minId = \min x : \{ \text{Actor}(x, y, _), y = \text{'John'} \}$

What does this even mean???

Can't use variable that are not aggregated in the outer /head atoms

CSE 414 - Spring 2018 36

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Counting

```
Q(c) :- c = count : { Actor(_, y, _), y = 'John' }
```

No variable here!

Meaning (in SQL, assuming no NULLs)

```
SELECT count(*) as c
FROM Actor as a
WHERE a.name = 'John'
```

CSE 414 - Spring 2018

37

Actor(id, fname, lname)
Casts(pid, mid)
Movie(id, name, year)

Grouping

```
Q(y,c) :- Movie(_,_,y), c = count : { Movie(x,_,y) }
```

Meaning (in SQL)

```
SELECT m.year, count(*)
FROM Movie as m
GROUP BY m.year
```

CSE 414 - Spring 2018

38

ParentChild(p,c)

Example

For each person, compute the total number of descendants

```
// for each person, compute his/her descendants
```

CSE 414 - Spring 2018

39

ParentChild(p,c)

Example

For each person, compute the total number of descendants

```
// for each person, compute his/her descendants
D(x,y) :- ParentChild(x,y).
```

CSE 414 - Spring 2018

40

ParentChild(p,c)

Example

For each person, compute the total number of descendants

```
// for each person, compute his/her descendants
D(x,y) :- ParentChild(x,y).
D(x,z) :- D(x,y), ParentChild(y,z).
```

CSE 414 - Spring 2018

41

ParentChild(p,c)

Example

For each person, compute the total number of descendants

```
// for each person, compute his/her descendants
D(x,y) :- ParentChild(x,y).
D(x,z) :- D(x,y), ParentChild(y,z).

// For each person, count the number of descendants
```

CSE 414 - Spring 2018

42

Example

For each person, compute the total number of descendants

```
// for each person, compute his/her descendants
D(x,y) :- ParentChild(x,y).
D(x,z) :- D(x,y), ParentChild(y,z).

// For each person, count the number of descendants
T(p,c) :- D(p,_), c = count : { D(p,y) }.
```

Example

For each person, compute the total number of descendants

```
// for each person, compute his/her descendants
D(x,y) :- ParentChild(x,y).
D(x,z) :- D(x,y), ParentChild(y,z).

// For each person, count the number of descendants
T(p,c) :- D(p,_), c = count : { D(p,y) }.

// Find the number of descendants of Alice
```

Example

For each person, compute the total number of descendants

```
// for each person, compute his/her descendants
D(x,y) :- ParentChild(x,y).
D(x,z) :- D(x,y), ParentChild(y,z).

// For each person, count the number of descendants
T(p,c) :- D(p,_), c = count : { D(p,y) }.

// Find the number of descendants of Alice
Q(d) :- T(p, d), p = "Alice".
```

Negation: use "!"

Find all descendants of Alice,
who are not descendants of Bob

```
// for each person, compute his/her descendants
D(x,y) :- ParentChild(x,y).
D(x,z) :- D(x,y), ParentChild(y,z).

// Compute the answer: notice the negation
Q(x) :- D("Alice",x), !D("Bob",x).
```