

Introduction to Database Systems

CSE 414

Lecture 12: NoSQL

Class Overview

- Unit 1: Intro
- Unit 2: Relational Data Models and Query Languages
- Unit 3: Non-relational data
 - NoSQL
 - Json
 - SQL++
- Unit 4: RDMBS internals and query optimization
- Unit 5: Parallel query processing
- Unit 6: DBMS usability, conceptual design
- Unit 7: Transactions
- Unit 8: Advanced topics (time permitting)

Two Classes of Database Applications

- OLTP (Online Transaction Processing)
 - Queries are simple lookups: 0 or 1 join
E.g., find customer by ID and their orders
 - Many updates. E.g., insert order, update payment
 - **Consistency** is critical: **transactions** (more later)
- OLAP (Online Analytical Processing)
 - aka “Decision Support”
 - Queries have many joins, and group-by’s
E.g., sum revenues by store, product, clerk, date
 - No updates

NoSQL Motivation


- Originally motivated by Web 2.0 applications
 - E.g. Facebook, Amazon, Instagram, etc
 - Web startups need to scaleup from 10 to 100000 users very quickly
- Needed: very large scale OLTP workloads
- Give up on consistency
- Give up OLAP

RDBMS Review: Serverless

Desktop



User




SQLite:

- One data file
- One user
- One DBMS application
- **Consistency** is easy
- But only a limited number of scenarios work with such model

Disk

RDBMS Review: Client-Server


H/w 3



Connection (JDBC, ODBC)

- One server running the database
- Many clients, connecting via the ODBC or JDBC (Java Database Connectivity) protocol

RDBMS Review: Client-Server



- One server running the database
- Many clients, connecting via the ODBC or JDBC (Java Database Connectivity) protocol

Client-Server

- One *server* that runs the DBMS (or RDBMS):
 - Your own desktop, or
 - Some beefy system, or
 - A cloud service (SQL Azure)


Client-Server

- One *server* that runs the DBMS (or RDBMS):
 - Your own desktop, or
 - Some beefy system, or
 - A cloud service (SQL Azure)
- Many *clients* run apps and connect to DBMS
 - Microsoft's Management Studio (for SQL Server), or
 - psql (for postgres)
 - Some Java program (HW8) or some C++ program

Client-Server

- One *server* that runs the DBMS (or RDBMS):
 - Your own desktop, or
 - Some beefy system, or
 - A cloud service (SQL Azure)
- Many *clients* run apps and connect to DBMS
 - Microsoft's Management Studio (for SQL Server), or
 - psql (for postgres)
 - Some Java program (HW8) or some C++ program
- Clients “talk” to server using JDBC/ODBC protocol


Web Apps: 3 Tier



Browser




Web Apps: 3 Tier





Web Apps: 3 Tier

Web-based applications




Connection
(e.g., JDBC)




Web Apps: 3 Tier



Web-based applications



Connection
(e.g., JDBC)







HTTP/SSL



Replicate
App server
for scaleup

OS: 3 Tier

Web-based applications



Why not replicate DB server?


HTTP/SSL





Replicate
App server
for scaleup

OS: 3 Tier


Web-based applications



Connection
(e.g., JDBC)



HTTP/SSL



Why not replicate DB server?
Consistency!




Replicating the Database

- Two basic approaches:
 - Scale up through **partitioning**
 - Scale up through **replication**
- **Consistency** is much harder to enforce


Scale Through Partitioning

- Partition the database across many machines in a cluster
 - Database now fits in main memory
 - Queries spread across these machines
- Can increase throughput
- Easy for writes but reads become expensive!



Scale Through Replication

- Create multiple copies of each database partition
- Spread queries across these replicas
- Can increase throughput and lower latency
- Can also improve fault-tolerance
- Easy for reads but writes become expensive!



Relational Model → NoSQL

- Relational DB: difficult to replicate/partition
- Given Supplier(sno,...), Part(pno,...), Supply(sno, pno)
 - Partition: we may be forced to join across servers
 - Replication: local copy has inconsistent versions
 - Consistency is hard in both cases (why?)
- NoSQL: simplified data model
 - Given up on functionality
 - Application must now handle joins and consistency

Data Models

Taxonomy based on data models:

- ☞ • **Key-value stores**
 - e.g., Project Voldemort, Memcached
- **Document stores**
 - e.g., SimpleDB, CouchDB, MongoDB

Key-Value Stores Features

- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)

Key-Value Stores Features


- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)
- **Operations**
 - `get(key)`, `put(key,value)`
 - Operations on value not supported

Key-Value Stores Features

- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)
- **Operations**
 - get(key), put(key,value)
 - Operations on value not supported
- **Distribution / Partitioning**

Aside: Hash Functions

- A function that maps any data to a “hash value” (e.g., an integer)



Aside: Hash Functions

- Example: data and hash value are integers
- Simple hash function:
 - $h(\text{key}) = \text{key \% 42};$
 - $h(10) = 10$
 - $h(2) = 2$
 - $h(50) = 8$
- What does this have to do with data distribution?

Key-Value Stores Features

- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)
- **Operations**
 - `get(key)`, `put(key,value)`
 - Operations on value not supported
- **Distribution / Partitioning** – w/ hash function
 - No replication: key k is stored at server $h(k)$
 - 3-way replication: key k stored at $h_1(k), h_2(k), h_3(k)$

How does `get(k)` work? How does `put(k,v)` work?

Flights(fid, date, carrier, flight_num, origin, dest, ...)

Carriers(cid, name)

Example

- How would you represent the Flights data as key, value pairs?

How does query processing work?

Flights(fid, date, carrier, flight_num, origin, dest, ...)

Carriers(cid, name)

Example

- How would you represent the Flights data as key, value pairs?
- Option 1: key=fid, value=entire flight record

How does query processing work?

Flights(fid, date, carrier, flight_num, origin, dest, ...)

Carriers(cid, name)

Example

- How would you represent the Flights data as key, value pairs?
- Option 1: key=fid, value=entire flight record
- Option 2: key=date, value=all flights that day

How does query processing work?

Flights(fid, date, carrier, flight_num, origin, dest, ...)

Carriers(cid, name)

Example

- How would you represent the Flights data as key, value pairs?
- Option 1: key=fid, value=entire flight record
- Option 2: key=date, value=all flights that day
- Option 3: key=(origin,dest), value=all flights between

How does query processing work?

Data Models

Taxonomy based on data models:

- **Key-value stores**
 - e.g., Project Voldemort, Memcached
- ☞ • **Document stores**
 - e.g., SimpleDB, CouchDB, MongoDB

Motivation


- In Key, Value stores, the Value is often a very complex object
 - Key = ‘2010/7/1’, Value = [all flights that date]
- Better: allow DBMS to understand the *value*
 - Represent *value* as a JSON (or XML...) document
 - [all flights on that date] = a JSON file
 - May search for all flights on a given date

Document Stores Features

- **Data model:** (key,document) pairs
 - Key = string/integer, unique for the entire data
 - Document = JSON, or XML
- **Operations**
 - Get/put document by key
 - Query language over JSON
- **Distribution / Partitioning**
 - Entire documents, as for key/value pairs

We will discuss JSON

Example: storing FB friends



As a graph

Add new attributes?

OR

Storing lists?

| Person1 | Person2 | is_friend |
|---------|---------|-----------|
| Peter | John | 1 |
| John | Mary | 0 |
| Mary | Phil | 1 |
| Phil | Peter | 1 |
| ... | ... | ... |

As a relation

We will learn the tradeoffs of different data models later this quarter

JSON

JSON - Overview

- JavaScript Object Notation = lightweight text-based open standard designed for human-readable data interchange. Interfaces in C, C++, Java, Python, Perl, etc.
- The filename extension is .json.

We will emphasize JSON as semi-structured data

JSON Syntax

```
{ "book": [  
    {"id": "01",  
     "language": "Java",  
     "author": "H. Javeson",  
     "year": 2015  
    },  
    {"id": "07",  
     "language": "C++",  
     "edition": "second"  
     "author": "E. Sepp",  
     "price": 22.25  
    }  
]
```

JSon vs Relational

- Relational data model
 - Rigid flat structure (tables)
 - Schema must be fixed in advanced
 - Binary representation: good for performance, bad for exchange
 - Query language based on Relational Calculus
- Semistructured data model / JSon
 - Flexible, nested structure (trees)
 - Does not require predefined schema ("self describing")
 - Text representation: good for exchange, bad for performance
 - Most common use: Language API; query languages emerging

JSon Terminology

- Data is represented in name/value pairs.
- Curly braces hold objects
 - Each object is a list of name/value pairs separated by ,(comma)
 - Each pair is a name followed by ':'(colon) followed by the value
- Square brackets hold arrays and values are separated by ,(comma).


JSon Data Structures

- Collections of name-value pairs:
 - {“**name1**”: value1, “**name2**”: value2, ...}
 - The “name” is also called a “key”
- Ordered lists of values:
 - [obj1, obj2, obj3, ...]

Avoid Using Duplicate Keys

The standard allows them, but many implementations don't

```
{"id": "07",  
 "title": "Databases",  
 "author": "Garcia-Molina",  
 "author": "Ullman",  
 "author": "Widom"}  
}
```



```
{"id": "07",  
 "title": "Databases",  
 "author": ["Garcia-Molina",  
 "Ullman",  
 "Widom"]}  
}
```


[]

JSon Datatypes

- Number
- String = double-quoted
- Boolean = true or false
- nullempty

JSon Semantics: a Tree !

```
{"person":  
  [ {"name": "Mary",  
      "address":  
        {"street": "Maple",  
         "no": 345,  
         "city": "Seattle"}},  
   {"name": "John",  
    "address": "Thailand",  
    "phone": 2345678}  
  ]  
}
```




JSon Data

- JSon is **self-describing**
- Schema elements become part of the data
 - Relational schema: `person(name,phone)`
 - In Json “`person`”, “`name`”, “`phone`” are part of the data, and are repeated many times
- Consequence: JSon is much more flexible
- JSon = **semistructured** data

Mapping Relational Data to JSON

Person

| name | phone |
|------|-------|
| John | 3634 |
| Sue | 6343 |
| Dirk | 6363 |



```
{"person":  
  [{"name": "John", "phone":3634},  
   {"name": "Sue", "phone":6343},  
   {"name": "Dirk", "phone":6383}  
 ]  
}
```

Mapping Relational Data to JSON

May inline foreign keys

Person

| name | phone |
|------|-------|
| John | 3634 |
| Sue | 6343 |

Orders

| personName | date | product |
|------------|------|---------|
| John | 2002 | Gizmo |
| John | 2004 | Gadget |
| Sue | 2002 | Gadget |

```
{"Person":  
  [{"name": "John",  
   "phone":3646,  
   "Orders": [{"date":2002,  
              "product":"Gizmo"},  
             {"date":2004,  
              "product":"Gadget"}]  
  },  
   {"name": "Sue",  
    "phone":6343,  
    "Orders": [{"date":2002,  
               "product":"Gadget"}]  
   }]  
}
```