

# Introduction to Database Systems CSE 414

## Lecture 22: Introduction to Transactions

CSE 414 - Spring 2018

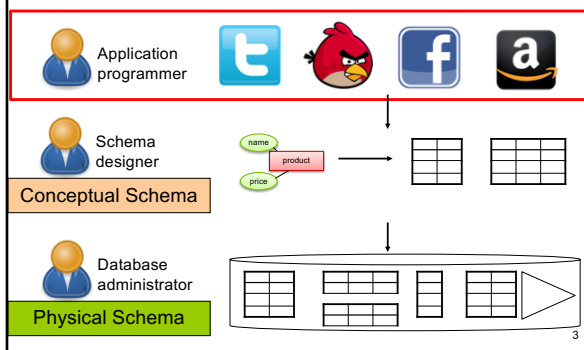
1

## Class Overview

- Unit 1: Intro
- Unit 2: Relational Data Models and Query Languages
- Unit 3: Non-relational data
- Unit 4: RDBMS internals and query optimization
- Unit 5: Parallel query processing
- Unit 6: DBMS usability, conceptual design
- Unit 7: Transactions
  - Locking and schedules
  - Writing DB applications
- Unit 8: Advanced topics

2

## Data Management Pipeline



3

## Transactions

- We use database transactions everyday
  - Bank \$\$\$ transfers
  - Online shopping
  - Signing up for classes
- For this class, a transaction is a series of DB queries
  - Read / Write / Update / Delete / Insert
  - Unit of work issued by a user that is independent from others

CSE 414 - Spring 2018

4

## What's the big deal?

CSE 414 - Spring 2018

5

## Challenges

- Want to execute many apps concurrently
  - All these apps read and write data to the same DB
- Simple solution: only serve one app at a time
  - What's the problem?
- Want: multiple operations to be executed atomically over the same DBMS

CSE 414 - Spring 2018

6

## What can go wrong?

- Manager: balance budgets among projects
  - Remove \$10k from project A
  - Add \$7k to project B
  - Add \$3k to project C
- CEO: check company's total balance
  - SELECT SUM(money) FROM budget;
- This is called a dirty / inconsistent read  
aka a **WRITE-READ** conflict

CSE 414 - Spring 2018

7

## What can go wrong?

- App 1:  
SELECT inventory FROM products WHERE pid = 1
- App 2:  
UPDATE products SET inventory = 0 WHERE pid = 1
- App 1:  
SELECT inventory \* price FROM products  
WHERE pid = 1
- This is known as an unrepeatable read  
aka **READ-WRITE** conflict

CSE 414 - Spring 2018

8

## What can go wrong?

Account 1 = \$100  
Account 2 = \$100  
Total = \$200

- |  |  |
|--|--|
| <ul style="list-style-type: none"> <li>• App 1:               <ul style="list-style-type: none"> <li>– Set Account 1 = \$200</li> <li>– Set Account 2 = \$0</li> </ul> </li> <li>• App 2:               <ul style="list-style-type: none"> <li>– Set Account 2 = \$200</li> <li>– Set Account 1 = \$0</li> </ul> </li> <li>• At the end:               <ul style="list-style-type: none"> <li>– Total = \$200</li> </ul> </li> </ul> | <ul style="list-style-type: none"> <li>• App 1: Set Account 1 = \$200</li> <li>• App 2: Set Account 2 = \$200</li> <li>• App 1: Set Account 2 = \$0</li> <li>• App 2: Set Account 1 = \$0</li> <li>• At the end:               <ul style="list-style-type: none"> <li>– Total = \$0</li> </ul> </li> </ul> |
|--|--|


This is called the lost update aka **WRITE-WRITE** conflict

CSE 414 - Spring 2018

9

## What can go wrong?

- Buying tickets to the next Bieber concert:
  - Fill up form with your mailing address
  - Put in debit card number
  - Click submit
  - Screen shows money deducted from your account
  - [Your browser crashes]



Lesson:  
Changes to the database  
should be **ALL or NOTHING**

CSE 414 - Spring 2018

10

## Transactions

- Collection of statements that are executed atomically (logically speaking)

```
BEGIN TRANSACTION
[SQL statements]
COMMIT or
ROLLBACK (=ABORT)
```

```
[single SQL statement]
```

If BEGIN... missing,  
then TXN consists  
of a single instruction

CSE 414 - Spring 2018

## Transactions Demo

CSE 414 - Spring 2018

12

## Turing Awards in Data Management



Charles Bachman, 1973  
*IDS and CODASYL*



Ted Codd, 1981  
*Relational model*



Jim Gray, 1998  
*Transaction processing*



Michael Stonebraker, 2014  
*INGRES and Postgres*



CSE 414 - Spring 2018

13

## Know your chemistry transactions: ACID

- **Atomic**
  - State shows either all the effects of txn, or none of them
- **Consistent**
  - Txn moves from a DBMS state where integrity holds, to another where integrity holds
    - remember integrity constraints?
- **Isolated**
  - Effect of txns is the same as txns running one after another (i.e., looks like batch mode)
- **Durable**
  - Once a txn has committed, its effects remain in the database

CSE 414 - Spring 2018

14

## Atomic

- **Definition:** A transaction is ATOMIC if all its updates must happen or not at all.
- **Example:** move \$100 from A to B
  - UPDATE accounts SET bal = bal - 100  
WHERE acct = A;
  - UPDATE accounts SET bal = bal + 100  
WHERE acct = B;
  - BEGIN TRANSACTION;
  - UPDATE accounts SET bal = bal - 100  
WHERE acct = A;
  - UPDATE accounts SET bal = bal + 100  
WHERE acct = B;
  - COMMIT;

CSE 414 - Spring 2018

15

## Isolated

- **Definition** An execution ensures that txns are isolated, if the effect of each txn is as if it were the only txn running on the system.

CSE 414 - Spring 2018

16

## Consistent

- Recall: integrity constraints govern how values in tables are related to each other
  - Can be enforced by the DBMS, or ensured by the app
- How consistency is achieved by the app:
  - App programmer ensures that txns only takes a consistent DB state to another consistent state
  - DB makes sure that txns are executed atomically
- Can defer checking the validity of constraints until the end of a transaction

CSE 414 - Spring 2018

17

## Durable

- A transaction is durable if its effects continue to exist after the transaction and even after the program has terminated
- How?
  - By writing to disk!
  - More in CSE 444

CSE 414 - Spring 2018

18

## Rollback transactions

- If the app gets to a state where it cannot complete the transaction successfully, execute ROLLBACK
- The DB returns to the state prior to the transaction
- What are examples of such program states?

CSE 414 - Spring 2018

19

## ACID

- Atomic
- Consistent
- Isolated
- Durable
- Enjoy this in HW8!
- Again: by default each statement is its own txn
  - Unless auto-commit is off then each statement starts a new txn

CSE 414 - Spring 2018

20

## Transaction Schedules

CSE 414 - Spring 2018

21

## Schedules

A **schedule** is a sequence of interleaved actions from all transactions

CSE 414 - Spring 2018

22

## Serial Schedule

- A *serial schedule* is one in which transactions are executed one after the other, in some sequential order
- **Fact:** nothing can go wrong if the system executes transactions serially
  - (up to what we have learned so far)
  - But DBMS don't do that because we want better overall system performance

CSE 414 - Spring 2018

23

## Example


A and B are elements in the database  
t and s are variables in txn source code

T1	T2
READ(A, t)	READ(A, s)
t := t+100	s := s*2
WRITE(A, t)	WRITE(A, s)
READ(B, t)	READ(B, s)
t := t+100	s := s*2
WRITE(B, t)	WRITE(B, s)

CSE 414 - Spring 2018

24


### Example of a (Serial) Schedule



CSE 414 - Spring 2018

25

### Another Serial Schedule



CSE 414 - Spring 2018

26

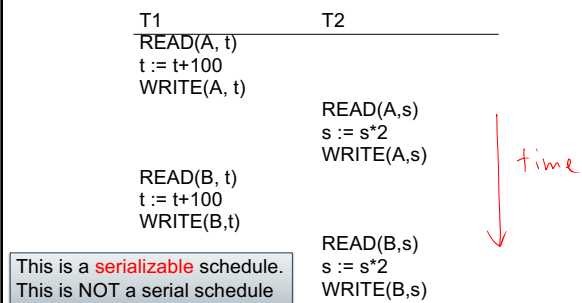
### Review: Serializable Schedule

A schedule is **serializable** if it is equivalent to a serial schedule

CSE 414 - Spring 2018

27

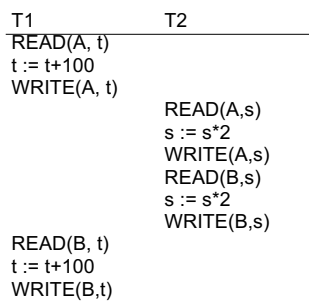
### A Serializable Schedule



CSE 414 - Spring 2018

28

### A Non-Serializable Schedule



CSE 414 - Spring 2018

29

### How do We Know if a Schedule is Serializable?

#### Notation:

T<sub>1</sub>: r<sub>1</sub>(A); w<sub>1</sub>(A); r<sub>1</sub>(B); w<sub>1</sub>(B)  
 T<sub>2</sub>: r<sub>2</sub>(A); w<sub>2</sub>(A); r<sub>2</sub>(B); w<sub>2</sub>(B)

Key Idea: Focus on *conflicting* operations

CSE 414 - Spring 2018

30